LaTeX Author Guidelines for CVPR Proceedings

Anonymous CVPR submission

Paper ID ****

Abstract

The ABSTRACT is to be in fully-justified italicized text, at the top of the left-hand column, below the author and affiliation information. Use the word "Abstract" as the title, in 12-point Times, boldface type, centered relative to the column, initially capitalized. The abstract is to be in 10-point, single-spaced type. Leave two blank lines after the Abstract, then begin the main text. Look at previous CVPR abstracts to get a feel for style and length.

1. Introduction

1.1. Context

Estimating depth from images is an important imaging problem, as dense depth maps are useful precursors for high-level scene understanding tasks like segmentation and pose estimation (cite cite cite) and mid-level perception (of which depth estimation is an important part) has been shown to be very useful for e.g. training robots to navigate their environments. While traditional approaches to depth estimation use multiple cameras or structure-from-motion, convolutional neural networks have also demonstrated reasonable performance on the so-called monocular depth estimation task, where the network is trained to produce a dense depth map given only a single RGB image of the scene.

1.2. Problem Statement

While deep monocular depth estimators have demonstrated strong performance (cite cite) and even some generalizability across scene types, the task they are solving is fundamentally underconstrained due to *inherent scale ambiguity*, i.e. the unresolvable tradeoff between size and distance in monocular images. This ambiguity could be resolved by adding an additional camera and calculating a disparity map, but this method still fails on textureless regions and areas with lots of occlusions. Other approaches use FMCW or time-of-flight LiDAR technologies, but these approaches are currently expensive and bulky.

1.3. Proposed approach

In this paper, we show that by augmenting the RGB image with a histogram of depth information, we can achieve substantially improved performance (and generalizability) over state-of-the-art monocular depth estimators. By performing an exact, weighted histogram matching on the output depth map of the depth estimator, we can match the depth histogram of the scene to the depth histogram of our estimate. Such a histogram can be captured relatively inexpensively using only a single pixel single-photon avalanche diode (SPAD) and pulsed laser illumination diffused over the field of view.

1.4. Impact

Our method is a lightweight postprocessing step that substantially improves the quality of depth maps produced by monocular depth estimators. It can be applied to any method to improve the accuracy instantly. (Our method also helps neural-network-based methods generalize across scene types easily.)

1.5. Limitations

Our method is not without limitations, however. It still requires a laser and single-pixel detector, and as such, is sensitive to ambient photons. Our method, which is a fundamentally a variant of histogram matching, is, like histogram matching, unable to transpose the values of pixels (i.e. if pixel a is farther than pixel b in the input, it will be farther than pixel b in the output). In other words, our method is not able to resolve ordinal depth errors (errors where an object is wrongly placed closer or farther relative to another object). Finally, our method is non-differentiable, and is therefore unsuitable for end-to-end optimization of multi-part networks.

- We introduce the idea of augmenting an RGB camera with a single-pixel SPAD to address scale ambiguity error in monocular depth estimators.
- We analyze our approach on indoor scenes using the NYU Depth v2 dataset. We demonstrate that our ap-

proach is able to resolve scale ambiguity while being fast and easy to implement.

- (Potentially) We investigate the ability of our method to help generalization of monocular depth estimators across scene types. We (hopefully) demonstrate that our method allows monocular depth estimators to perform well even on completely different scene types.
- We build a hardware prototype and evaluate the efficacy of our approach on real-world data.

2. Related Work

Depth Imaging

- stereo and multiview
- structured illumination and random patterns (kinect, etc.), active stereo
- time of flight (continuous wave and pulsed)
- what we do: like pulsed but much simpler setup; no scanning, no spad array, ...

Monocular Depth Estimation

- summary of architectures and cost functions: u-net type architecture with reverse huber loss
- what we do: same thing, but augment with global hints (inspired by these approaches, we do ...)

Deep Sensor Fusion global hints for super-resolution, colorization, depth estimation

- colorization
- david's 2018 paper for depth estimation and denoising (see david's 2019 sig paper for related work)
- what we do: slightly different application

Histogram Matching Histogram matching as an image processing technique

- Exact histogram matching paper used in this work
- Wasserstein-based optimization techniques for histogram-based regularization

3. Method

In this section, we describe the measurement model for a single-pixel time-of-flight lidar sensor under diffuse, pulsed laser illumination.

3.1. Measurement Model

Consider a laser which emits a pulse at time $t = 0$ with time-varying intensity $g(t)$ uniformly illuminating some 3D scene. We parameterize the geometry of the scene as a height map $z(x, y)$. Neglecting albedo and falloff effects, an ideal detector counting photon events from a location (x, y) in the time interval $(n\Delta t, (n + 1)\Delta t)$ would record

$$\lambda_{x,y}[n] = \int_{n\Delta t}^{(n+1)\Delta t} (f * g)(t - 2z(x, y)/c)dt \quad (1)$$

where c is the speed of light, and f is a function that models the temporal uncertainty in the detector. Single-photon avalanche diodes (SPADs) are highly sensitive photodetectors which are able to record single photon events with high temporal precision [**?**]. Since the detection of each photon can be described with a Bernoulli random variable, the total number of accumulated photons in this time interval follows a Poisson distribution according to

$$h[n] \sim \mathcal{P}\left(\sum_{x,y} \alpha_{x,y}\eta\lambda_{x,y}[n] + b\right) \quad (2)$$

where $\alpha_{x,y} = r_{x,y}/z(x, y)^2$ captures the attenuation of the photon counts due to the reflectance $r(x, y)$ of the scene and due to the inverse square falloff $1/z(x, y)^2$. In addition, η is the detection probability of a photon triggering a SPAD event, and $b = \eta a + d$ is the average number of background detections resulting from ambient photons a and erroneous "dark count" events d resulting from noise within the SPAD.

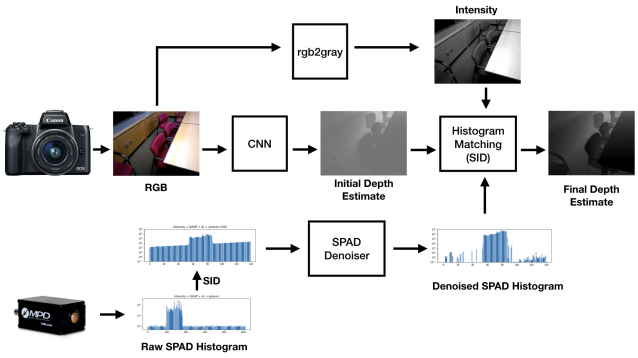

Figure 1: **Overview of the full pipeline** We use a CNN to get an initial per-pixel depth estimate. We then perform gradient descent to optimize that estimate using the SPAD forward model and the dual-Sinkhorn distance .

3.2. Monocular depth estimation with global depth hints

Given a single RGB image $I(x, y)$ and a vector of photon arrivals $h[n]$ described by equation 2, we seek to reconstruct the ground truth depth map $z(x, y)$. Our method has two parts. First, we initialize our estimate of the depth map from the single RGB image via a monocular depth estimator described below. Second, we refine this depth map using the captured measurements $h[n]$ via a process we call Differentiable Histogram Matching (DHM). Differentiable histogram matching is a tool for post-processing the image to match the depth map to the statistics we capture from the SPAD.

Initialization via CNN Convolutional Neural Networks have become increasingly capable of leveraging monocular depth cues to produce accurate estimates of depth from only a single image. We therefore choose to initialize our depth map estimate $\hat{z}^{(0)}(x, y)$ using a CNN. However, any depth estimator reliant on only a single view will be unable to resolve the inherent scale ambiguity in the scene resulting from the tradeoff between size of and distance to an object. The next step, differentiable histogram matching, will resolve this ambiguity using the depth information present in the SPAD histogram.

SPAD Denoising

- Discuss MLE for SPAD denoising
- Write optimization problem for SPAD denoising
- Show performance on a few examples

Exact Histogram Matching An image's *histogram* is a pair of vectors (h, b) where h_i is the number of pixels of the image whose value lies in the range $[b_i, b_i + 1)$. Then, given a source image S with histogram (h_s, b) and a target histogram (h_t, b), histogram matching generates a new image M such that $h_m \approx h_t$ and the pixel values in M are in the same relative order as in S.

3.3. Implementation Details

For the Monocular Depth Estimator, we use pretrained versions of the the Deep Ordinal Regression Network (DORN) [] and the DenseDepth Network. The exact histogram matching method is as described in [].

4. Simulation

4.1. Implementation Details

- Number of bins used, depth range, laser parameters, use of intensity image.
- Using

NYU Depth v2 The NYU Depth v2 Dataset consists of 249 training and 215 testing scenes of RGB-D data captured using a Microsoft Kinect. We used a version of DORN pretrained according to [**?**] as our CNN.

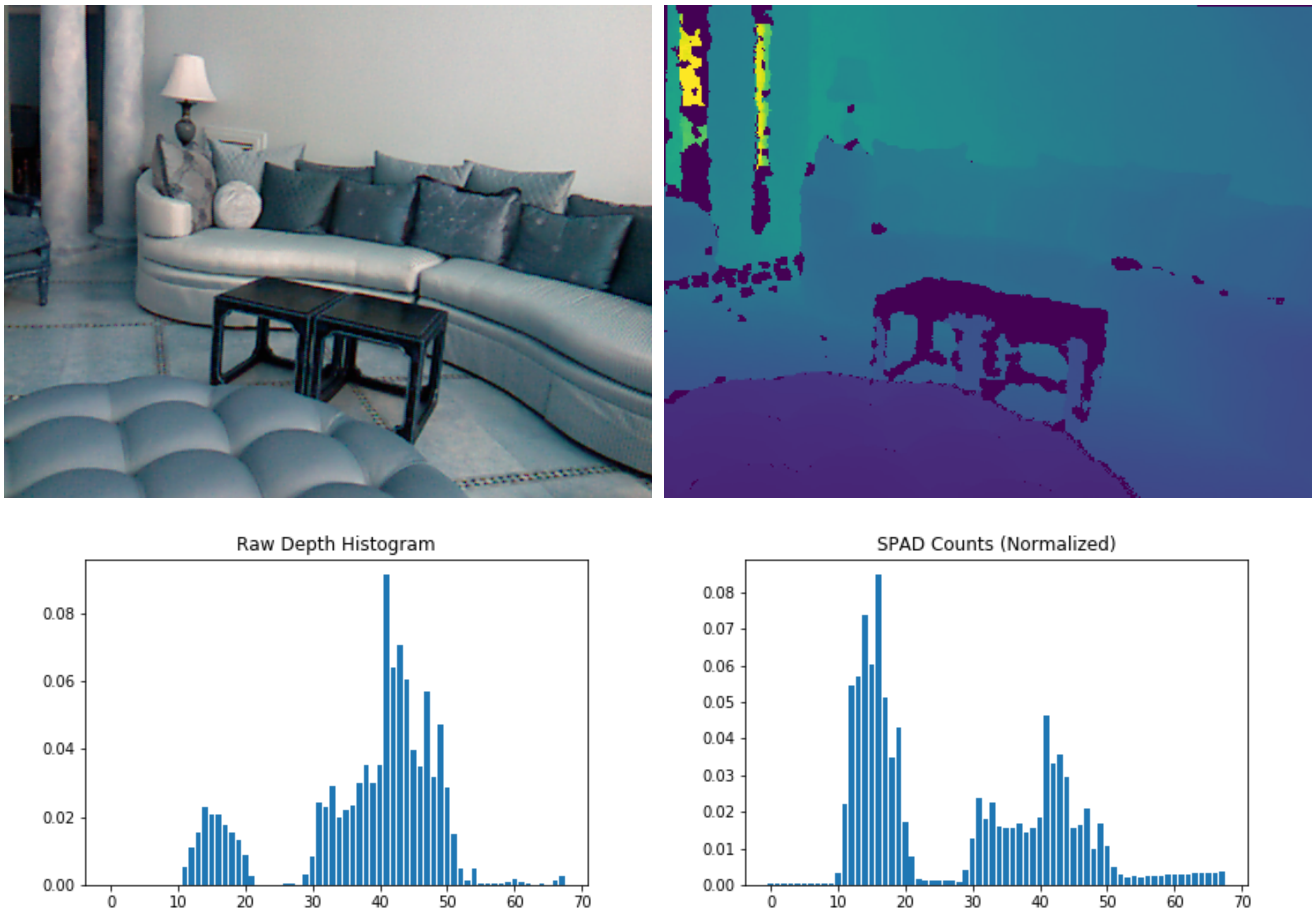

Table 1: Sample Image. Top Left is the RGB image. Top Right is ground truth depth. Bottom Left is Raw ground truth depth histogram. Bottom Right is simulated SPAD measurements. Notice how closer depths are magnified and far depths are attenuated.

5. Hardware Prototype

5.1. Setup

- Description of hardware used
- Images of scenes used

6. Discussion

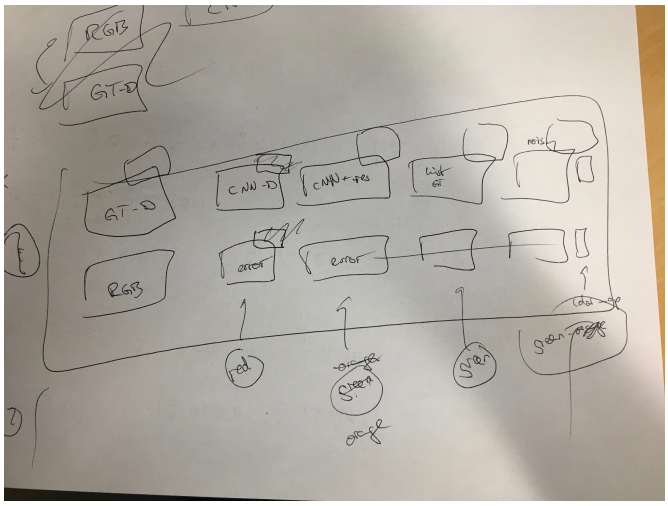

Figure 2: **Comparing our results with other methods**

	$\delta^1 \uparrow$	$\delta^2 \uparrow$	$\delta^3 \uparrow$	RMSE $\downarrow$	rel $\downarrow$	$log_{10} \downarrow$
Eigen et. al.	0.769	0.950	0.988	0.641	0.158	-
Laina et. al.	0.811	0.953	0.988	0.573	0.127	0.055
DORN	0.818	0.950	0.982	0.620	0.137	0.063
DORN (rescaled)	0.872	0.967	0.989	0.548	0.111	0.048
Alhashim, Wonka (2019)	0.847	0.973	0.994	0.548(0.461)	0.123	0.053
Alhashim, Wonka (2019) rescaled using GT depth"	0.888	**0.978**	**0.995**	**0.499**(0.409)	**0.106**	**0.045**
Ours (raw depth counts)	**0.899**	0.970	0.990	0.529	0.199	0.055
Ours (DORN) (intensity/falloff)	0.835	0.953	0.984	0.521	0.129	0.060
Ours (DenseDepth) (intensity/falloff)	0.867	0.974	0.994	0.445	0.114	0.050

Table 2: Results on the NYU Depth v2 test set [**?**].

model	hyperparams	$\delta 1 \uparrow$	$\delta 2 \uparrow$	$\delta 3 \uparrow$	$rel \downarrow$	$rmse \downarrow$	$log10 \downarrow$
dorn	CNN	0.846	0.954	0.983	0.120	0.501	0.053
	CNN + median rescaling	0.871	0.964	0.988	0.111	0.473	0.048
	CNN + GT histogram matching	0.906	0.972	0.990	0.095	0.419	0.040
	Ours (SBR=10)	0.903	0.970	0.989	0.091	0.422	0.040
	Ours (SBR=50)	0.906	0.971	0.990	0.089	0.410	0.039
	Ours (SBR=100)	0.906	0.971	0.990	0.090	0.408	0.039
densedepth	CNN	0.847	0.973	0.994	0.123	0.461	0.053
	CNN + median rescaling	0.888	0.978	0.995	0.106	0.409	0.045
	CNN + GT histogram matching	0.930	0.984	0.995	0.079	0.338	0.034
	Ours (SBR=10)	0.922	0.982	0.994	0.082	0.361	0.036
	Ours (SBR=50)	0.925	0.983	0.995	0.081	0.348	0.035
	Ours (SBR=100)	0.926	0.983	0.995	0.081	0.346	0.035

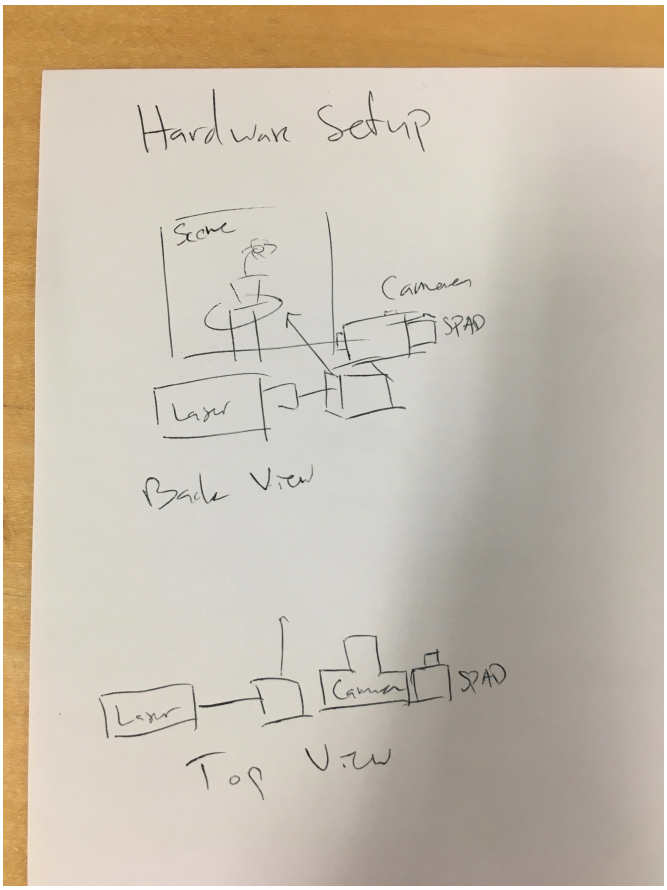

Figure 3: **Hardware setup**

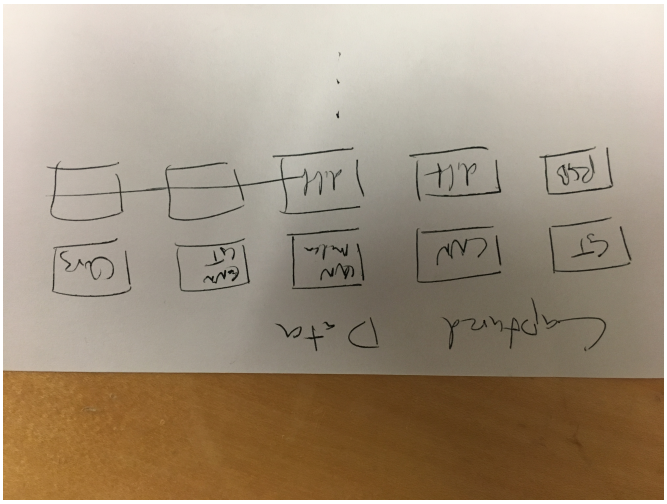

Figure 4: **Hardware results**